



# Operating System Concepts

Che-Wei Chang

[chewei@mail.cgu.edu.tw](mailto:chewei@mail.cgu.edu.tw)

Department of Computer Science and Information Engineering, Chang Gung University

# Contents

1. Introduction
2. System Structures
- ➔ 3. Process Concept
4. Multithreaded Programming
5. Process Scheduling
6. Synchronization
7. Deadlocks
8. Memory-Management Strategies
9. Virtual-Memory Management
10. File System
11. Implementing File Systems
12. Secondary-Storage Systems





# Chapter 3. Process Concept


# Objectives

- ▶ To introduce the notion of a process
- ▶ To describe the various features of processes, including scheduling, creation and termination, and communication
- ▶ To explore inter-process communication
- ▶ To describe communication in client-server systems




# Basic Process Concept

- ▶ A program is a **passive** entity stored on disk, and a process is an **active** entity
  - A program becomes process when the executable file is loaded into memory
  - The execution of a program started via GUI mouse clicks, the command line entry of its name, etc.
  - One program can be executed as several processes
- ▶ An operating system can execute a variety of programs
  - In batch systems: jobs
  - In time-shared systems: user programs or tasks



# Process in Memory

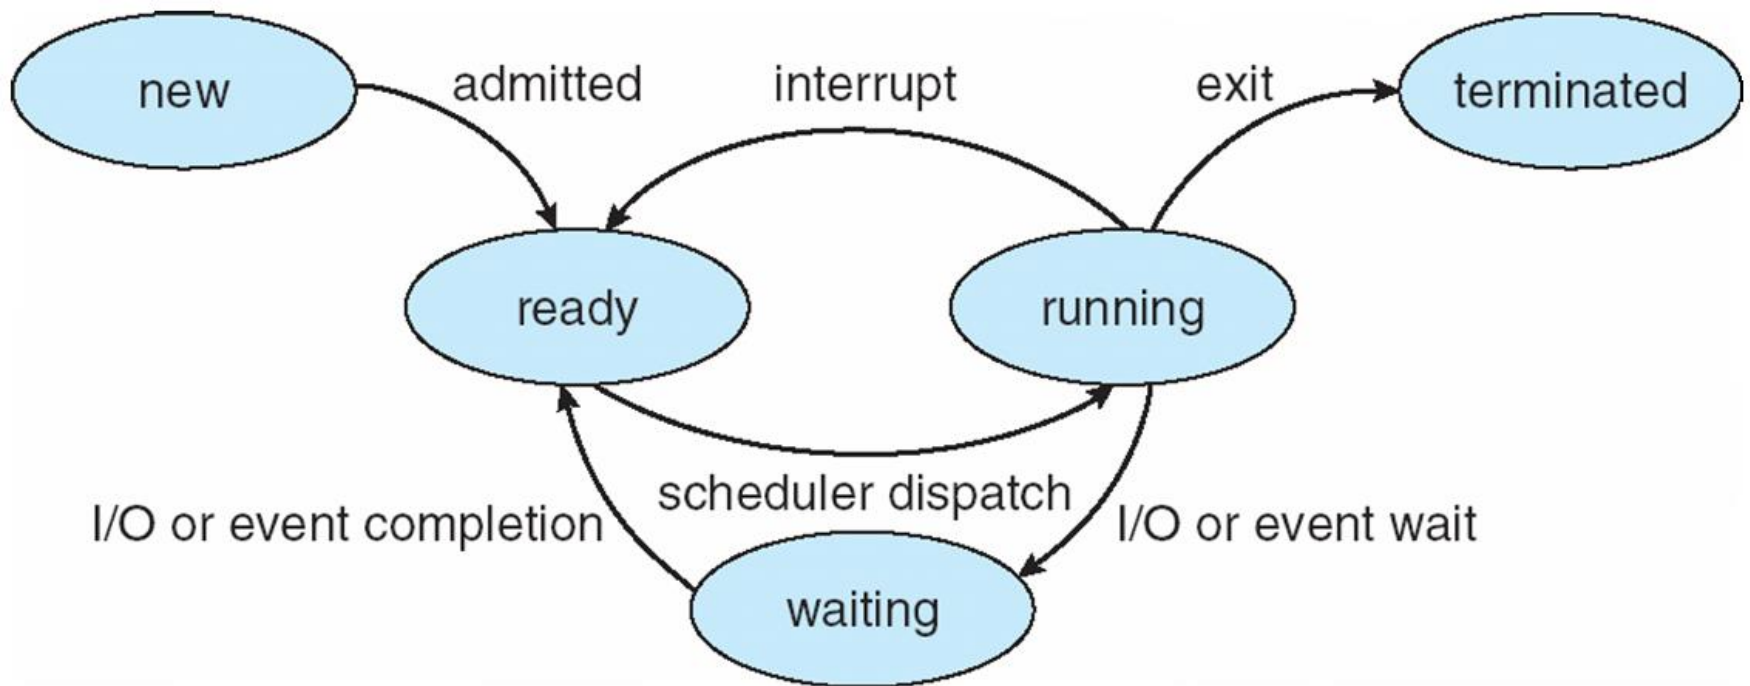


# Process States

- ▶ **New**: The process is being created
- ▶ **Running**: Instructions are being executed
- ▶ **Waiting**: The process is waiting for some event to occur
- ▶ **Ready**: The process is waiting to be assigned to a processor
- ▶ **Terminated**: The process has finished execution



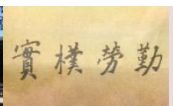
# Diagram of Process States






# Process Control Block (PCB)

- ▶ PCB: The repository for any information that may vary from process to process
  - Process state— running, waiting, etc
  - Program counter— location of the currently executed instruction
  - CPU registers— contents of all process-centric registers
  - CPU scheduling information— priorities, scheduling queue pointers
  - Memory-management information— memory allocated to the process
  - Accounting information— CPU used, clock time elapsed since start, time limits
  - I/O status information— I/O devices allocated to process, list of opened files

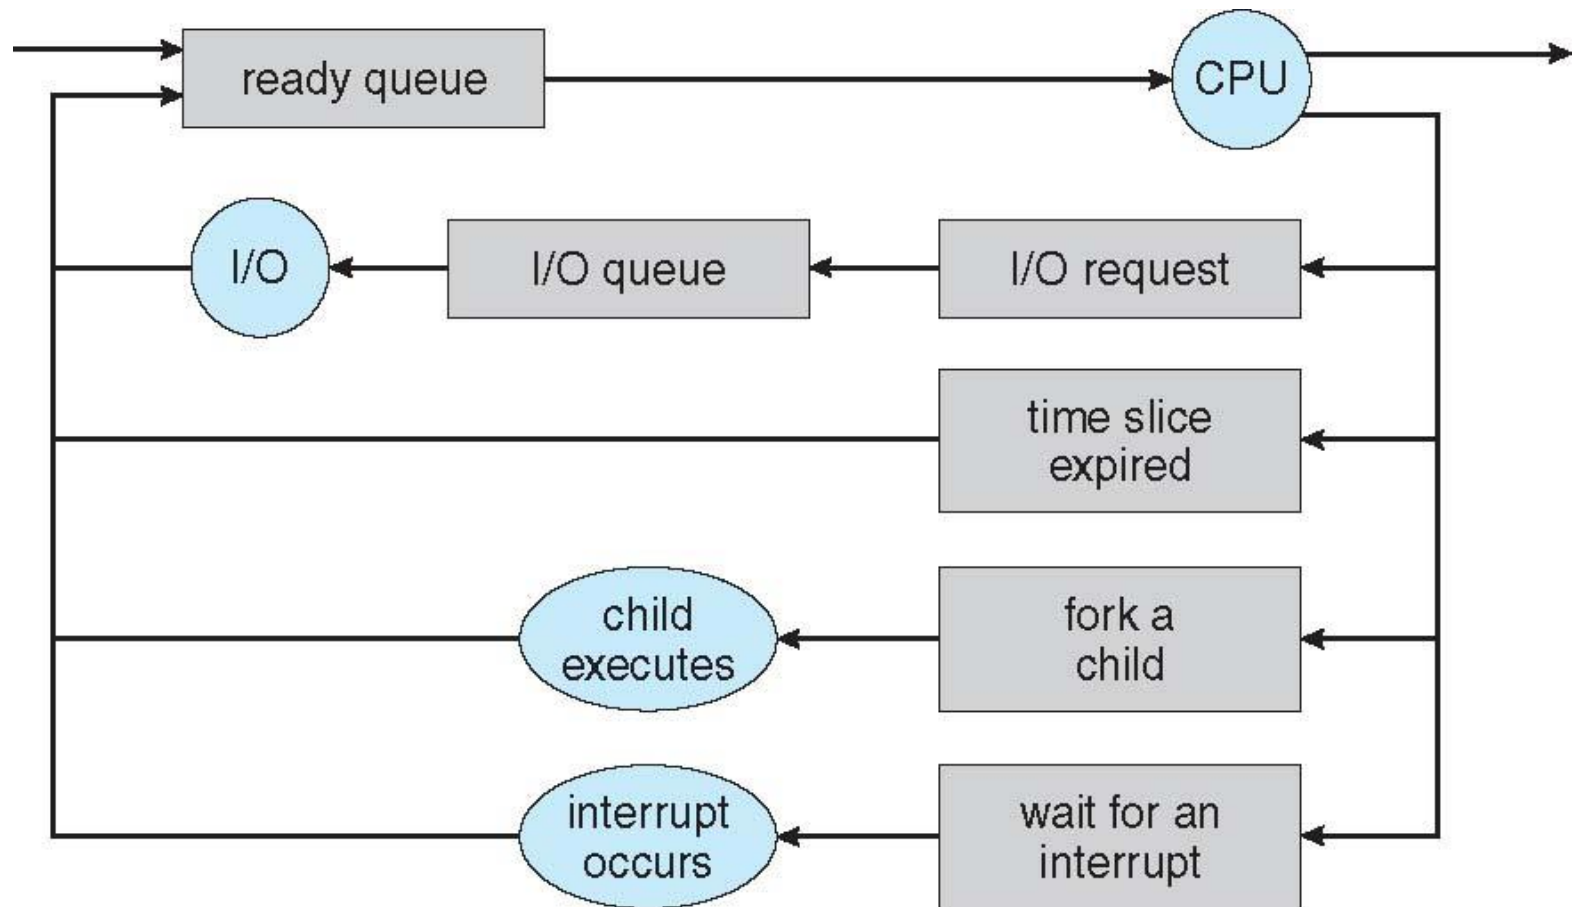


# Process Scheduling with PCB

- ▶ The goal of multiprogramming
  - Maximize CPU/resource utilization
- ▶ The goal of time sharing
  - Allow each user to interact with his/her program



# Process Scheduling- A Queueing Diagram




# Processor Schedulers

- ▶ **Long-term scheduler (or job scheduler)**— selects which processes should be brought into the ready queue
- ▶ **Short-term scheduler (or CPU scheduler)**— selects which process should be executed next and allocates CPU
- ▶ **Medium-term scheduler** can be added as **swapper**

# Long-Term Scheduler

- ▶ Processes can be described as either:
  - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
  - **CPU-bound process** – spends more time doing computations; few very long CPU bursts

- ▶ Long-term scheduler strives for good **process mix**



- ▶ Remarks :
  - Control the degree of multiprogramming
  - Can take more time in selecting processes because of a longer interval between executions
  - May not exist physically



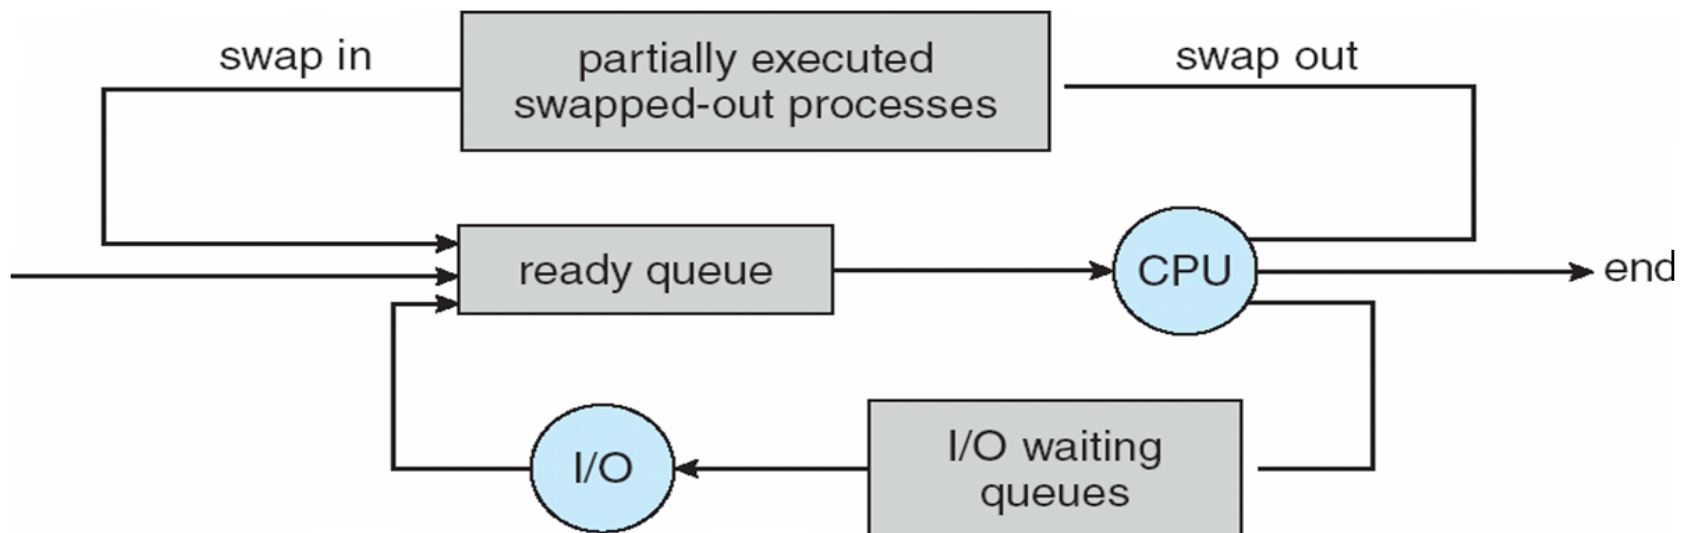
# Short-Term Scheduler

- ▶ Goal: To efficiently allocate the CPU to one of the ready processes according to some criteria
- ▶ Short-term scheduler is invoked very frequently (milliseconds) → must be fast
- ▶ In Linux, after version 2.6.23, the scheduler is the Completely Fair Scheduler (CFS)



# Medium-Term Scheduler

- ▶ Goal: Remove process from memory, store on disk, bring back in from disk to continue execution: it is also called “swapping”



# Process Scheduling– Context Switches

## ▶ Context Switch: Pure Overheads

- Save the state of the old process and load the state of the newly scheduled process.
  - The context of a process is usually reflected in PCB

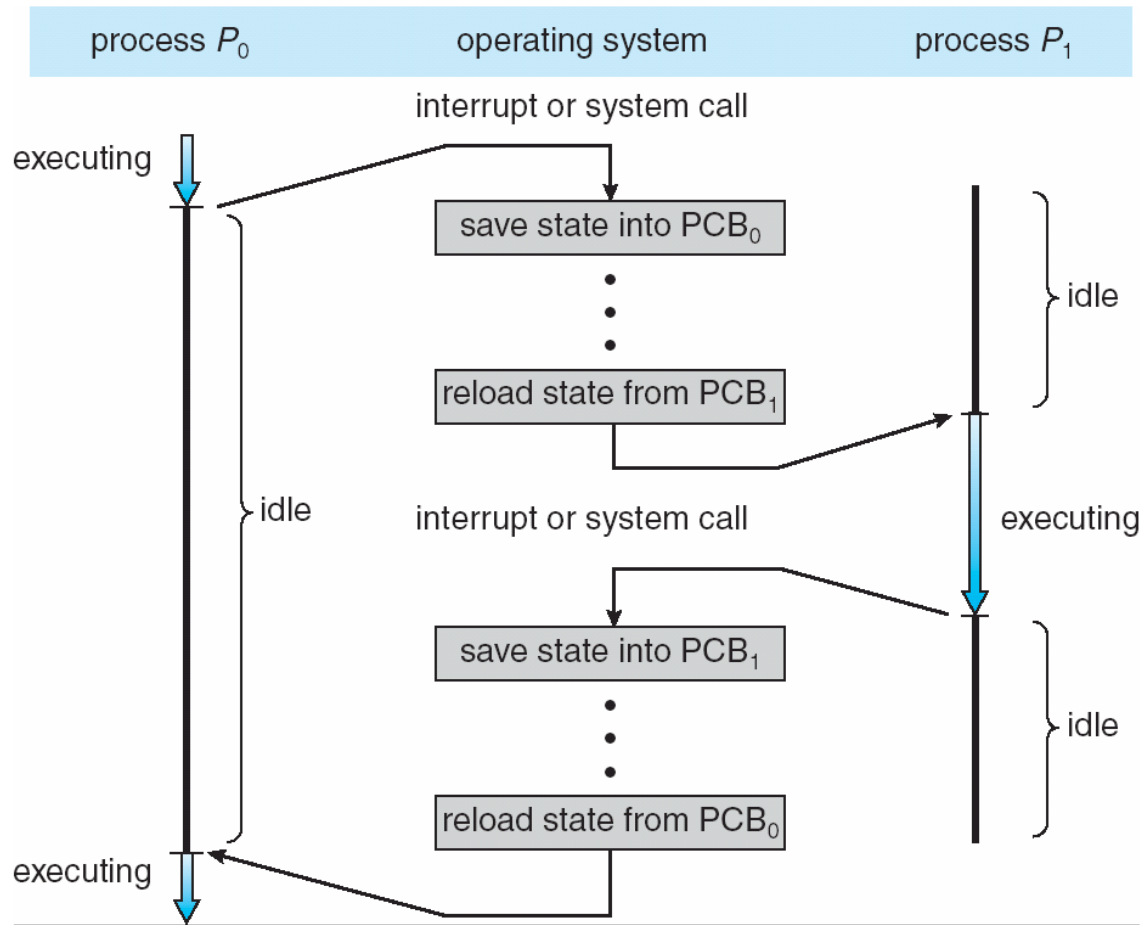
## ▶ Issues:

- The cost depends on the hardware support
  - e.g. processors with multiple register sets or computers with advanced memory management
- Threads, i.e., light-weight process (LWP), are introduced to break this bottleneck






# CPU Switch from Process to Process




# Parent and Child Processes

- ▶ Parent processes create child processes, which in turn create other processes, forming a tree of processes
- ▶ Generally, process identified and managed via a process identifier (PID)



# Process Creation

- ▶ Address Space
  - Child duplicate of parent
  - Child has a program loaded into it
- ▶ UNIX Examples
  - **fork()** system call creates new process
  - **exec()** system call used after a **fork()** to replace the process' memory space with a new program



# Process Termination

- ▶ Process executes last statement and asks the operating system to delete it: **exit()**
  - Wait the output data from child to parent: **wait()**
- ▶ Parent may terminate the execution of child processes: **abort()**
  - ➔ Child has exceeded allocated resources
  - ➔ Task assigned to child is no longer required
  - Receive the return value from child
  - Some operating systems do not allow child to continue if its parent terminates
    - All children should be terminated - **cascading termination**

# C Program Forking a Process

```
#include <sys/types.h>
#include <stdio.h>
#include <unistd.h>

int main()
{
    pid_t pid;

    /* fork a child process */
    pid = fork();

    if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        return 1;
    }
    else if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", NULL);
    }
    else { /* parent process */
        /* parent will wait for the child to complete */
        wait(NULL);
        printf("Child Complete");
    }

    return 0;
}
```



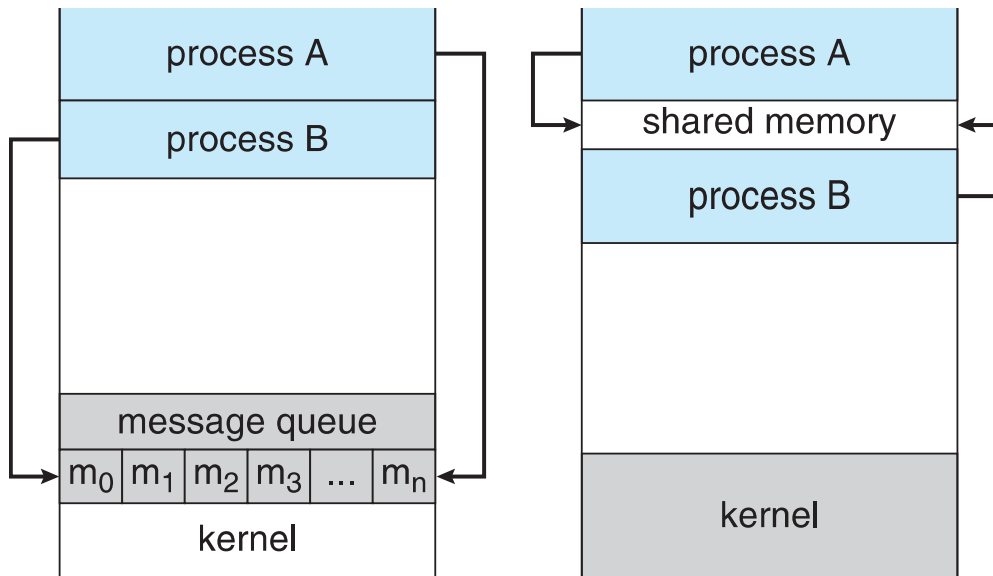
# Inter-Process Communication

- ▶ Processes within a system may be independent or cooperating
- ▶ Cooperating process can affect or be affected by other processes, including sharing data
- ▶ Reasons for cooperating processes:
  - Information sharing
  - Computation speedup
  - Modularity
  - Convenience
- ▶ Cooperating processes need inter-process communication (IPC)




# Two Models of IPC

- ▶ Shared Memory
  - Max Speed & Communication Convenience
- ▶ Message Passing
  - No Access Conflict & Easy Implementation



# Shared Memory IPC

- ▶ A Consumer-Producer Example:
  - Bounded buffer or unbounded buffer
    - Supported by inter-process communication (IPC) or by hand coding



`buffer[0...n-1]`

Initially,  $in=out=0$



# Shared Memory– Consumer

```
while (true)
{
    while (in == out);
    /* do nothing and have to wait */
    next_consumed = buffer[out];
    out = (out + 1) % BUFFER_SIZE;
    ... /* use the consumed item */
}
```



# Shared Memory– Producer

```
while (true)
{
    ... /* produce a new item */
    while (((in + 1) % BUFFER SIZE) == out) ;
    /* do nothing */
    buffer[in] = next_produced;
    in = (in + 1) % BUFFER SIZE;
}
```



# Message Passing IPC

- ▶ Logical Implementation of Message Passing
  - Fixed/variable message size
  - Symmetric/asymmetric communication
  - Direct/indirect communication
  - Synchronous/asynchronous communication
  - Automatic/explicit buffering
  - Send by copy or reference



# Direct Message Passing

- ▶ Processes must name each other explicitly:
  - **send** ( $P$ , *message*) – send a message to process  $P$
  - **receive**( $Q$ , *message*) – receive a message from process  $Q$
- ▶ Properties of the communication link
  - Links are established automatically
  - A link is associated with exactly one pair of communicating processes
  - Between each pair there exists exactly one link
  - The link may be unidirectional, but is usually bi-directional



# Indirect Communication


- ▶ Messages are directed and received from mailboxes (also referred to as ports)
  - **send**( $A, message$ ) – send a message to mailbox  $A$
  - **receive**( $A, message$ ) – receive a message from mailbox  $A$
- ▶ Properties of the communication link
  - Links are established only if processes share a common mailbox
  - A link may be associated with many processes
  - Each pair of processes may share several communication links
  - Links may be unidirectional or bi-directional



# Issues of Indirect Communication

## ► Mailbox sharing

- $P_1$ ,  $P_2$ , and  $P_3$  share mailbox A
- $P_1$ , sends;  $P_2$  and  $P_3$  receive
- Who gets the message?



## ► Solutions

- Allow a link to be associated with at most two processes
- Allow only one process at a time to execute a receive operation
- Allow the system to select arbitrarily the receiver



# IPC Synchronization

- ▶ Synchronous Message Passing IPC
  - Blocking send has the sender block until the message is received
  - Blocking receive has the receiver block until a message is available
- ▶ Asynchronous Message Passing IPC
  - Non-blocking send has the sender send the message and continue
  - Non-blocking receive has the receiver receive a valid message or null



# IPC Buffering

- ▶ The capacity of a link: the number of messages could be held in the link
  - Zero capacity – 0 messages
    - Sender must wait for receiver
  - Bounded capacity – finite length of  $n$  messages
    - Sender must wait if link is full
  - Unbounded capacity – infinite length
    - Sender never waits
- ▶ The last two items are for asynchronous communication and may need **acknowledgement**





# Examples of IPC Systems – POSIX

## ► POSIX Shared Memory

- Process first creates shared memory segment

```
shm_fd = shm_open(name, O_CREAT | O_RDWR,  
0666);
```

- Set the size of the object

```
ftruncate(shm_fd, 4096);
```

- Memory map the object

```
ptr = mmap(0, 4096, PROT_WRITE, MAP_SHARED,  
shm_fd, 0);
```

- Now the process could write to the shared memory

```
sprintf(ptr, "Writing to shared memory");
```



# Examples of IPC Systems – Mach

- ▶ Mach – A message-based OS from the Carnegie Mellon University
  - When a task is created, two special mailboxes, called ports, are also created.
    - The *Kernel* mailbox is used by the kernel to communicate with the tasks
    - The *Notify* mailbox is used by the kernel sends notification of event occurrences.



# Three IPC System Calls on Mach


- ▶ `msg_send`
  - Options when mailbox is full:
    - Wait indefinitely
    - Return immediately
    - Wait at most for n ms
    - Temporarily cache a message: only one message to a full mailbox can be pending at any time for a sending thread
- ▶ `msg_receive`
  - Only one task can own & have a receiving privilege of a mailbox
  - Options when mailbox is empty:
    - Wait indefinitely
    - Return immediately
    - Wait at most for n ms
- ▶ `msg_rpc`
  - Remote Procedure Calls



# Communication in Client-Server Systems

## ▶ Socket

- An endpoint for communication identified by an IP address concatenated with a port number
  - A client-server architecture
- ▶ /etc/services: 23-telnet, 21-ftp, 80-web server, etc.

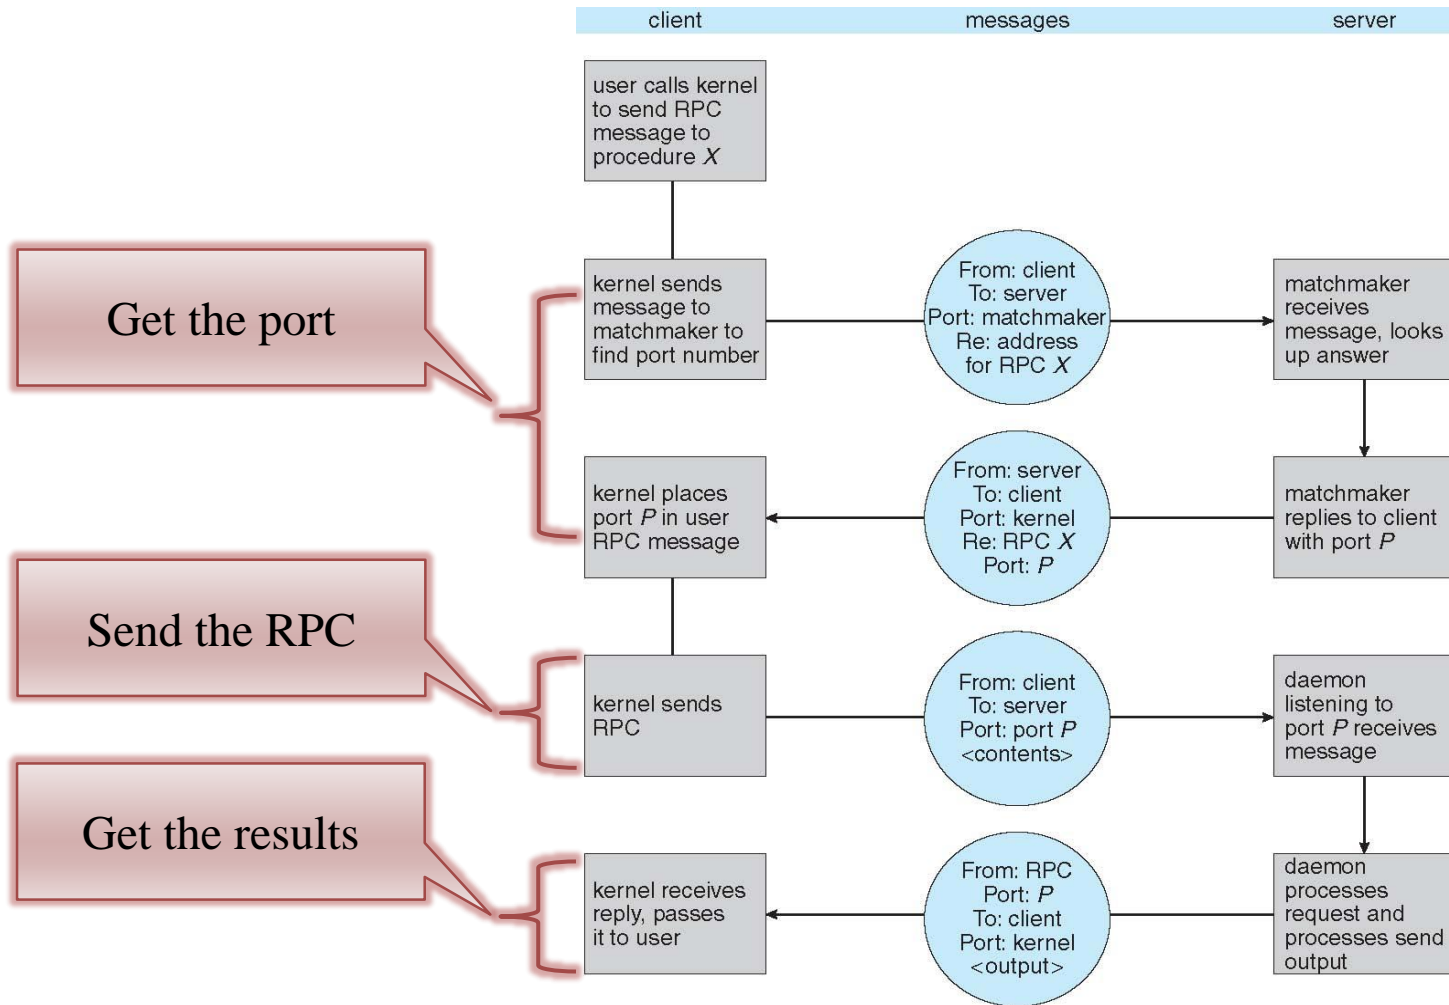


# Remote Procedure Calls

- ▶ A way to abstract the procedure-call mechanism for use between systems with network connection
- ▶ Stubs at the client site
  - One for each RPC
  - Locate the proper port and marshal parameters
- ▶ Stubs at the server site
  - Receive the message
  - Invoke the procedure and return the results
- ▶ Data representation handled via the External Data Representation (XDL) format to account for different architectures
  - Big-endian and little-endian



# Execution of RPC




# Pipes

- ▶ Acts as a conduit allowing two processes to communicate
- ▶ Issues
  - Is communication unidirectional or bidirectional?
  - In the case of two-way communication, is it half-duplex or full-duplex?
  - Must there exist a relationship (i.e. parent-child) between the communicating processes?
  - Can the pipes be used over a network?



# Ordinary Pipes

- ▶ Ordinary Pipes allow communication in the standard producer-consumer style
- ▶ Producer writes to the **write-end** of the pipe
- ▶ Consumer reads from the **read-end** of the pipe
- ▶ Ordinary pipes are therefore unidirectional
- ▶ Require parent-child relationship between communicating processes





# Named Pipes

- ▶ Named Pipes are more powerful than ordinary pipes
- ▶ Communication is bidirectional
- ▶ No parent-child relationship is necessary between the communicating processes
- ▶ Several processes can use the named pipe for communication
- ▶ Provided on both UNIX and Windows systems

