



Operating System Concepts

Che-Wei Chang

chewei@mail.cgu.edu.tw

Department of Computer Science and Information Engineering, Chang Gung University

Contents

1. Introduction
2. System Structures
3. Process Concept
4. Multithreaded Programming
- ➔ 5. Process Scheduling
6. Synchronization
7. Deadlocks
8. Memory-Management Strategies
9. Virtual-Memory Management
10. File System
11. Implementing File Systems
12. Secondary-Storage Systems

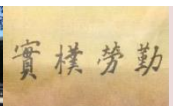




Chapter 5. Process Scheduling

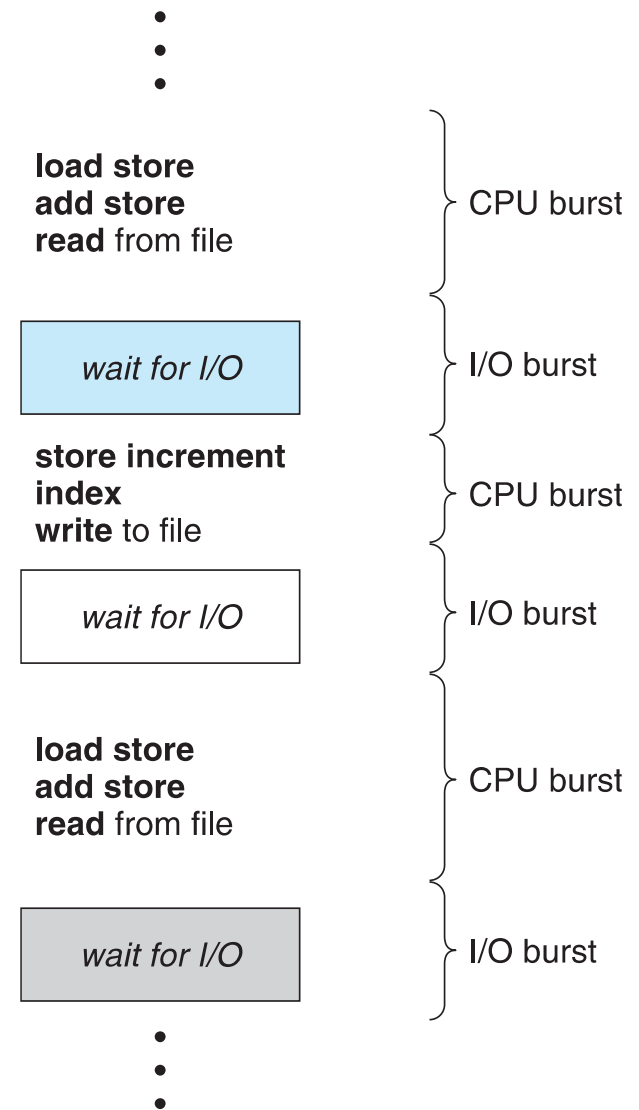
Objectives

- ▶ To introduce CPU scheduling, which is the basis for multi-programmed operating systems
- ▶ To describe various CPU-scheduling algorithms
- ▶ To discuss evaluation criteria for selecting a CPU scheduling algorithm for a particular system
- ▶ To examine the scheduling algorithms of several operating systems




Basic Concepts

- ▶ CPU–I/O Burst Cycle
 - Process execution consists of a cycle of CPU execution and I/O waiting
- ▶ Process Execution
 - CPU-bound programs tend to have a few very long CPU bursts
 - IO-bound programs tend to have many very short CPU bursts



Histogram of CPU-burst Times

- ▶ The distribution can help in selecting an appropriate CPU scheduling algorithms



CPU Scheduler

- ▶ Short-term scheduler selects a process among the processes in the ready queue, and allocates the CPU to the selected process
 - Queue may be ordered in various ways
- ▶ CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
- ▶ Scheduling under 1 and 4 is nonpreemptive
- ▶ All other scheduling is preemptive



Dispatcher

- ▶ Dispatcher module gives control of the CPU to the process selected by the short-term scheduler
 - switching context
 - switching to user mode
 - jumping to the proper location in the user program to resume that process
- ▶ Dispatch latency – the time it takes for the dispatcher to stop one process and start another running



Scheduling Criteria

- ▶ Why?
 - Different scheduling algorithms may favor one class of processes over another
- ▶ Criteria
 - CPU Utilization
 - Throughput
 - Turnaround Time: $(\text{Completion Time}) - (\text{Start Time})$
 - Waiting Time: Waiting in the Ready Queue
 - Response Time: First Response Time



Scheduling Algorithms

- ▶ First-Come, First-Served Scheduling (FIFO)
- ▶ Shortest-Job-First Scheduling (SJF)
- ▶ Priority Scheduling
- ▶ Round-Robin Scheduling (RR)
- ▶ Multilevel Queue Scheduling
- ▶ Multilevel Feedback Queue Scheduling
- ▶ Multiple-Processor Scheduling



First-Come, First-Served (FCFS) Scheduling

- ▶ The process which requests the CPU first is allocated the CPU
- ▶ Properties:
 - Non-preemptive scheduling
 - CPU might be hold for an extended period



A Scheduling Example of FCFS (1 / 2)

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- ▶ Suppose that the processes arrive in the order: P_1 , P_2 , P_3



- ▶ Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- ▶ Average waiting time: $(0 + 24 + 27)/3 = 17$



A Scheduling Example of FCFS (2/2)

- ▶ Suppose that the processes arrive in the order:
 - P_2, P_3, P_1
- ▶ The Gantt chart for the schedule is:



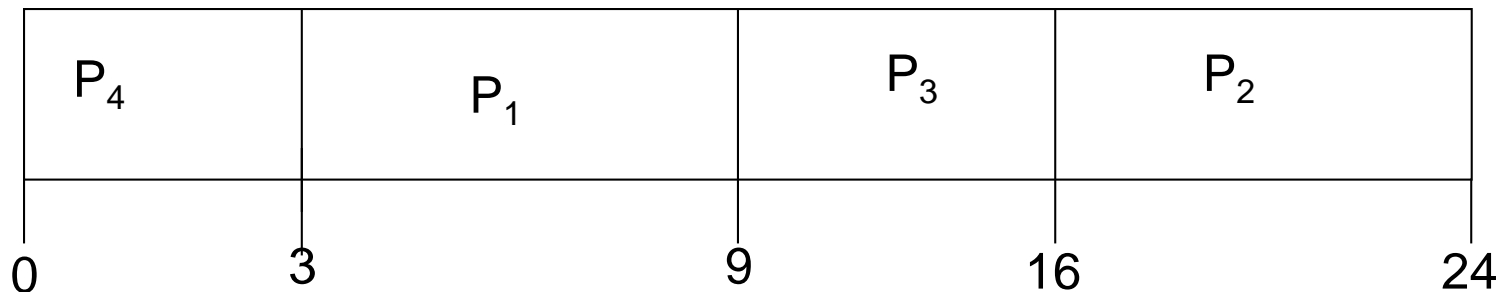
- ▶ Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- ▶ Average waiting time: $(6 + 0 + 3)/3 = 3$
- ▶ **Convoy effect** – short processes behind long a process



Shortest-Job-First (SJF) Scheduling

<u>Process</u>	<u>Burst Time</u>
P_1	6
P_2	8
P_3	7
P_4	3

► SJF scheduling chart



► Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$



SJF Scheduling Analysis


- ▶ Non-preemptive SJF scheduling is optimal when processes are all ready at time 0
 - The minimum average waiting time
- ▶ It is difficult to know the length of the next CPU request
 - Prediction of the next CPU burst time using exponential averaging
 1. t_n = actual length of n^{th} CPU burst
 2. τ_{n+1} = predicted value for the next CPU burst
 3. $\alpha, 0 \leq \alpha \leq 1$
 4. Define : $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$




Preemptive SJF Scheduling

► Preemptive or Non-preemptive?

- Criteria such as AWT (Average Waiting Time)



or



Shortest-Remaining-Time-First Scheduling



Priority Scheduling


- ▶ A priority number (integer) is associated with each process
- ▶ The CPU is allocated to the process with the highest priority
- ▶ Priority Assignment
 - Internally defined – use some measurable quantity, such as the number of open files, $\frac{\text{Average CPU Burst}}{\text{Average I/O Burst}}$
 - Externally defined – set by criteria external to the OS, such as the criticality levels of jobs



A Scheduling Example with Priority Scheduling

<u>Process</u>	<u>CPU Burst Time</u>	<u>Priority</u>
P1	10	3
P2	1	1
P3	2	3
P4	1	4
P5	5	2

Gantt Graph



$$\text{Average waiting time} = (6+0+16+18+1)/5 = 8.2$$



Issues of Priority Scheduling

- ▶ Problem: Starvation – low priority processes may never execute
- ▶ Solution: Aging – as time progresses increase the priority of the process
- ▶ A Special Case: SJF is priority scheduling where priority is the inverse of predicted next CPU burst time



Round Robin (RR) Scheduling


- ▶ Each process gets a small unit of CPU time (time quantum)
- ▶ After this time has elapsed, the process is preempted and added to the end of the ready queue
- ▶ If there are n processes in the ready queue and the time quantum is q
 - Each process gets $1/n$ of the CPU time in chunks of at most q time units at once
 - No process waits more than $(n-1)q$ time units



A Scheduling Example of RR Scheduling

<u>Process</u>	<u>CPU Burst Time</u>
P1	24
P2	3
P3	3

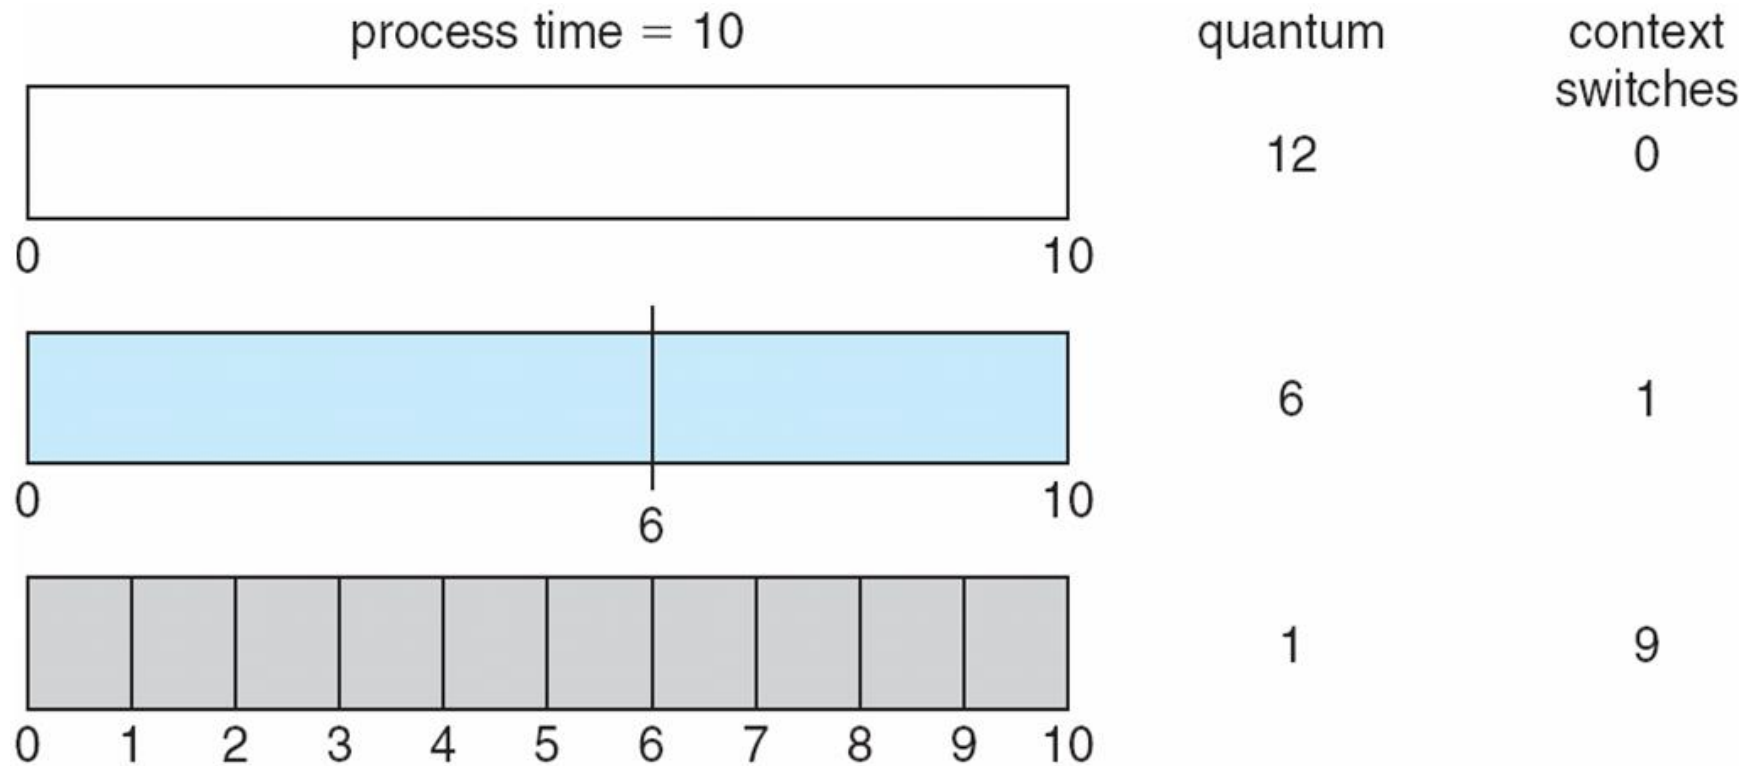
Time slice = 4



$$AWT = ((10-4) + (4-0) + (7-0))/3 = 17/3 = 5.66$$




Time Quantum and Context Switch



Issues of RR Scheduling

- ▶ Time quantum too large → FIFO
- ▶ Time quantum too small → Time quantum must be large with respect to context switch time, otherwise overhead is too high
 - Time quantum usually 10 ms to 100ms
 - Context switch $< 10 \mu\text{s}$
- ▶ A rule of thumb is that 80 percent of the CPU bursts should be shorter than the time quantum




Issues of RR Scheduling — Turnaround Time

process (10ms)


quantum = 10

quantum = 1


P1



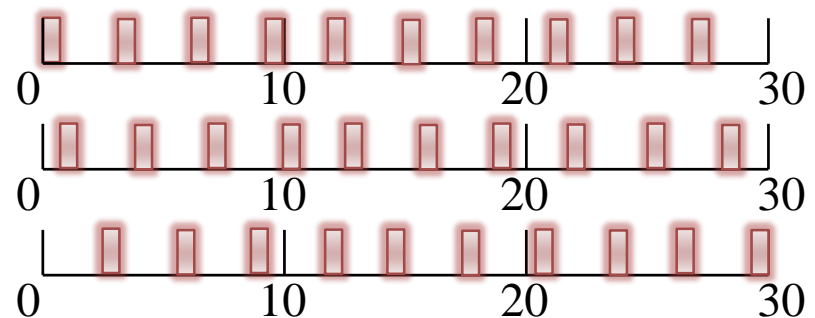
P2



P3



Average Turnaround Time
 $= (10+20+30)/3 = 20$



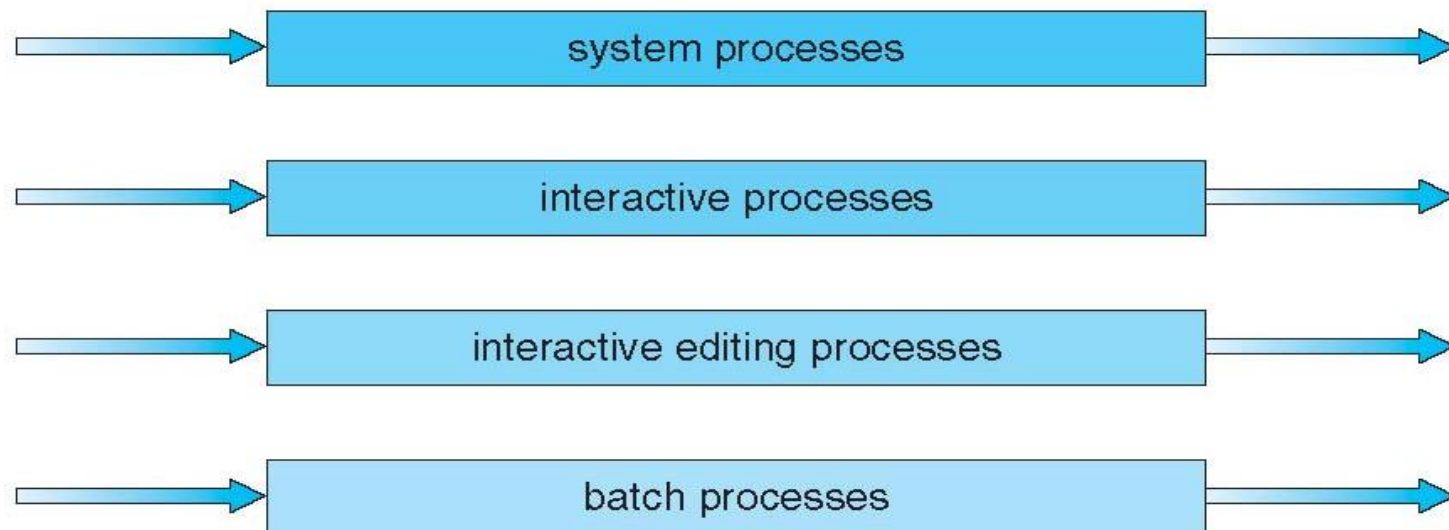
$ATT = (28+29+30)/3 = 29$

$\Rightarrow 80\% \text{ CPU Burst} < \text{time slice}$



Multilevel Queue Scheduling

- ▶ Partition the ready queue into several separate queues
 - ➔ Processes can be classified into different groups and permanently assigned to one queue



Multilevel Queue Scheduling

- ▶ Intra-queue scheduling
 - Independent choice of scheduling algorithms
 - e. g., foreground – RR, and background – FCFS
- ▶ Inter-queue scheduling
 - Fixed-priority preemptive scheduling
 - e.g., foreground queues always have absolute priority over the background queues
 - Time slice between queues
 - e.g., 80% CPU is given to foreground processes, and 20% CPU to background processes



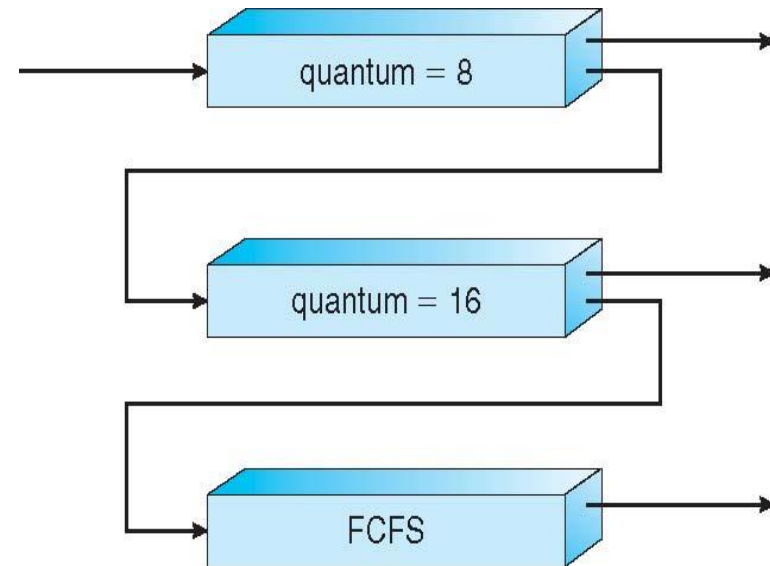
Multilevel Feedback Queue Scheduling

- ▶ A process can move between the various queues
 - Aging can be implemented this way
- ▶ Multilevel-feedback-queue scheduler defined by the following parameters:
 - number of queues
 - scheduling algorithms for each queue
 - method used to determine when to upgrade a process
 - method used to determine when to demote a process
 - The method to determine which queue a newly ready process will enter



An Example of Multilevel Feedback Queue

- ▶ Three queues:
 - Q_0 – RR with time quantum 8 milliseconds
 - Q_1 – RR time quantum 16 milliseconds
 - Q_2 – FCFS
- ▶ Scheduling
 - Do jobs in Q_0 first and then Q_1 and then Q_2
 - A new job enters queue Q_0
 - When it gains CPU, job receives 8 milliseconds
 - If it does not finish in 8 milliseconds, job is moved to queue Q_1
 - At Q_1 each job receives 16 additional milliseconds
 - If it still does not complete, it is preempted and moved to queue Q_2



Thread Scheduling

- ▶ To run on a CPU, user threads must be mapped to an associated kernel thread
- ▶ Local Scheduling
 - Contention Scope: Process-Contention Scope (PCS)
 - How the threads library decides which thread to put onto an available kernel thread
- ▶ Global Scheduling
 - Contention Scope: System-Contention Scope (SCS)
 - How the kernel decides which kernel thread to run on CPU next



Multiple-Processor Scheduling

- ▶ CPU scheduling in a system with multiple CPUs
- ▶ A Homogeneous System
 - Processors are identical in terms of their functionality
- ▶ A Heterogeneous System
 - Programs must be compiled for instructions on proper processors




Homogeneous Processors

- ▶ **Asymmetric multiprocessing** – only one processor accesses the system data structures, alleviating the need for data sharing
- ▶ **Symmetric multiprocessing (SMP)** – each processor is self-scheduling, all processes in common ready queue, or each processor has its own private queue of ready processes



Multiple-Processor Scheduling—Processor Affinity

- ▶ A process might prefer to run on specific processors
 - Hard affinity: `sched_setaffinity()`
 - Soft affinity: non-uniform memory access



Multiple-Processor Scheduling— Load Balancing

- ▶ Attempt to keep the workload evenly distributed across all processors in an SMP system
- ▶ Push migration
 - A specific task periodically checks the load on each processor and evenly distributes the load by moving processes from overloaded to idle or less-busy processors
- ▶ Pull migration
 - An idle processor pulls a waiting task from a busy processor

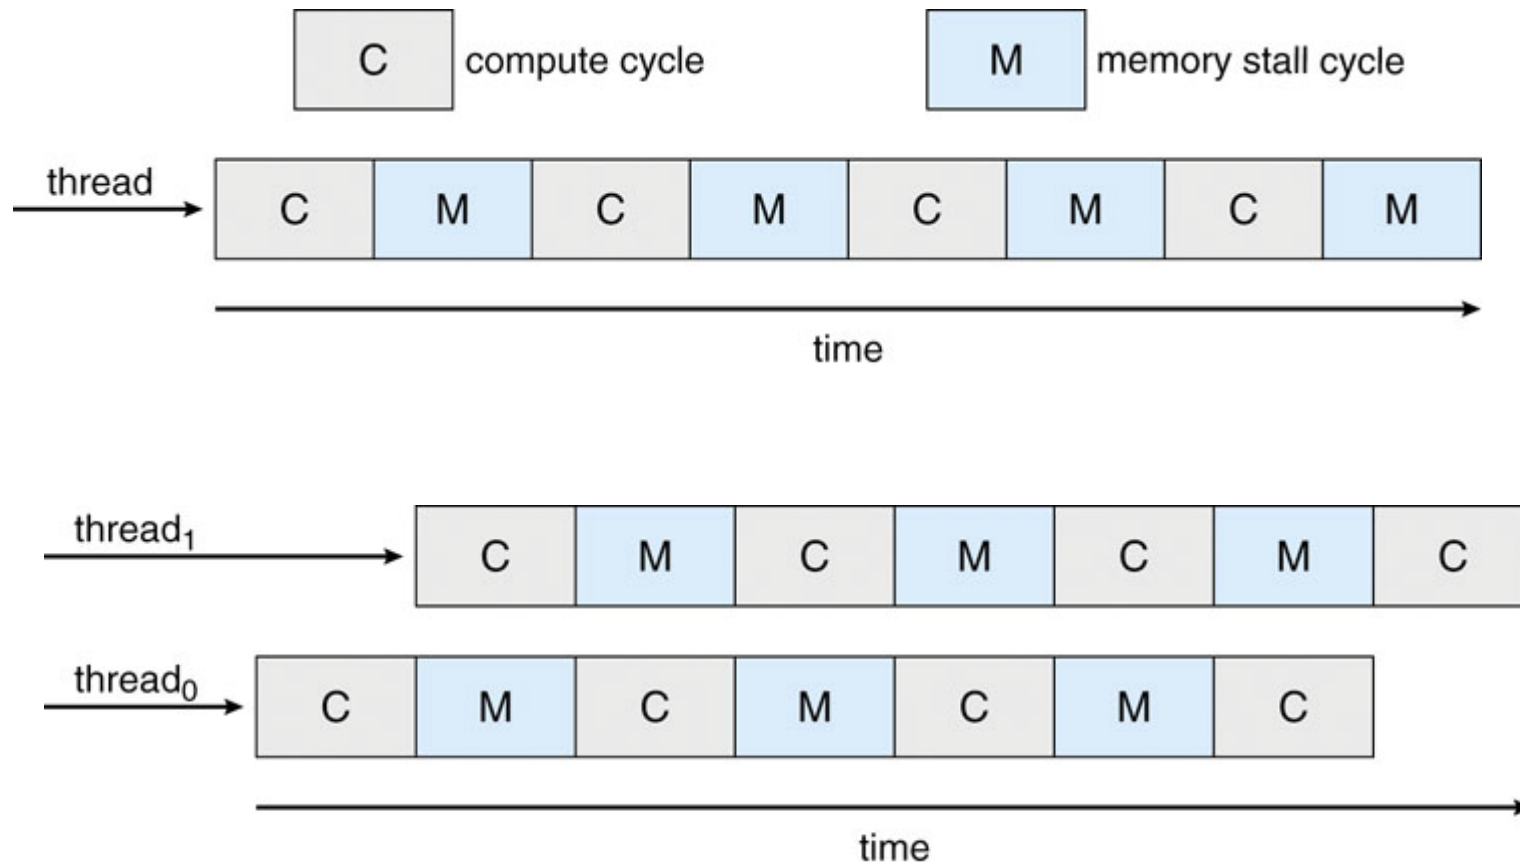


Multicore Processors

- ▶ Multicore Processor: A physical chip with multiple processor cores.
 - ▶ Scheduling Issues:
 - Memory Stall
 - Coarse-Grained Multithreading
 - Thread execution until a long latency
 - Fine-Grained Multithreading
 - Better architecture design for switching
- Multiple Hardware Threads



Multithreaded Multicore (Hyper-Threading) System

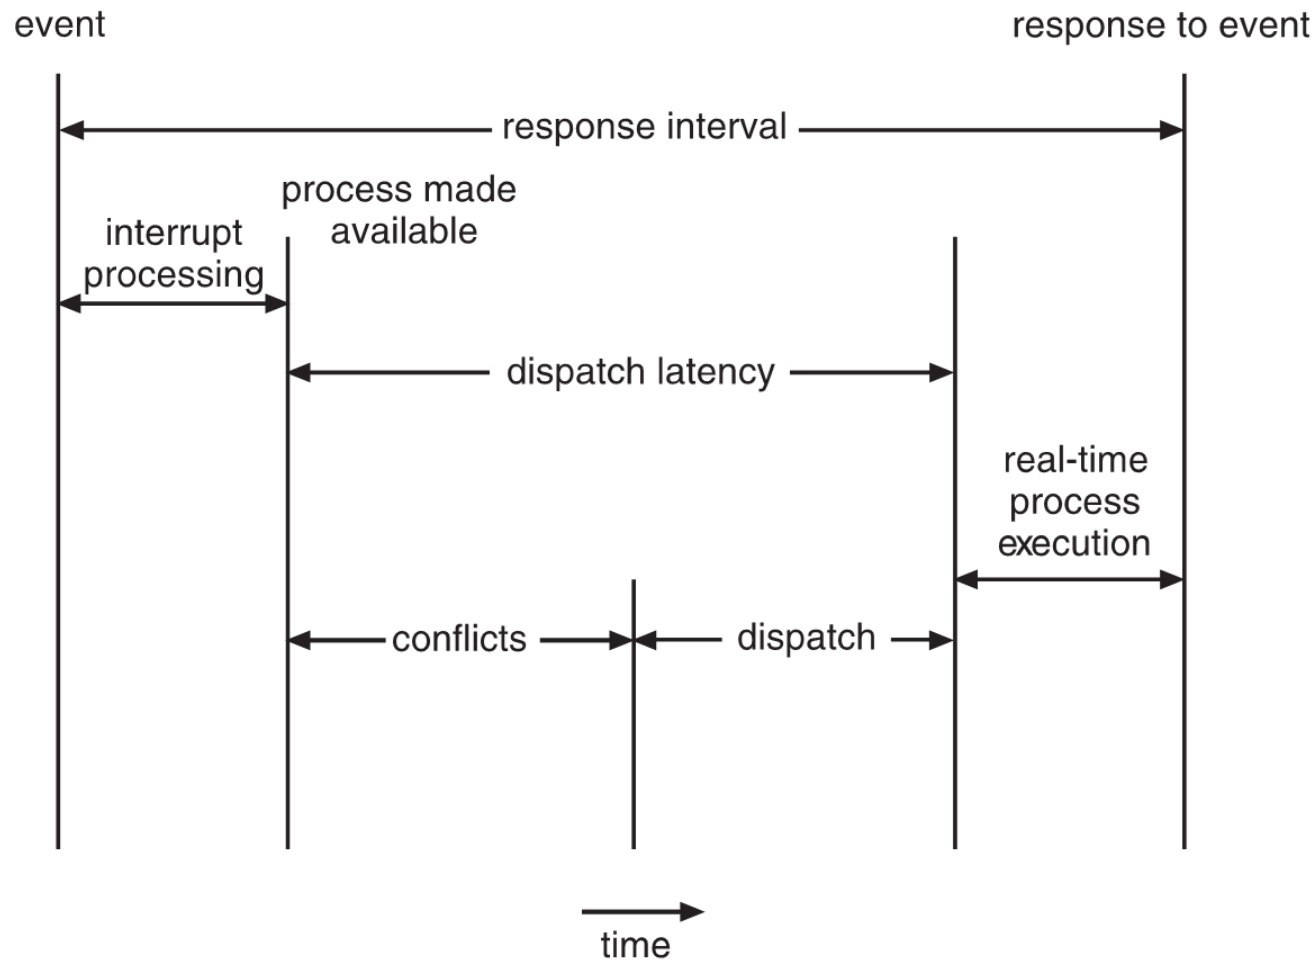


Real-Time Scheduling (1 / 2)

- ▶ Each task (process) has to be completed before its **deadline**
- ▶ **Soft real-time systems** – try to serve a real-time task by its deadline
- ▶ **Hard real-time systems** – a real-time task must be served by its deadline
- ▶ Two types of latencies affect performance
 1. Interrupt latency – time from arrival of interrupt to start of routine that serves the interrupt
 2. Dispatch latency – time for schedule to take current process off CPU and switch to another



Real-Time Scheduling (2/2)



Operating System Examples – Linux in Version 2.6.23 + (1 / 3)

► Completely Fair Scheduler (CFS)

- CFS scheduler maintains per task **virtual run time** in variable **vruntime**
- Associated with decay factor based on priority of task:
 - lower priority → higher decay rate
- Normal default priority yields virtual run time = actual run time
- To decide next task to run, scheduler picks task with lowest virtual run time

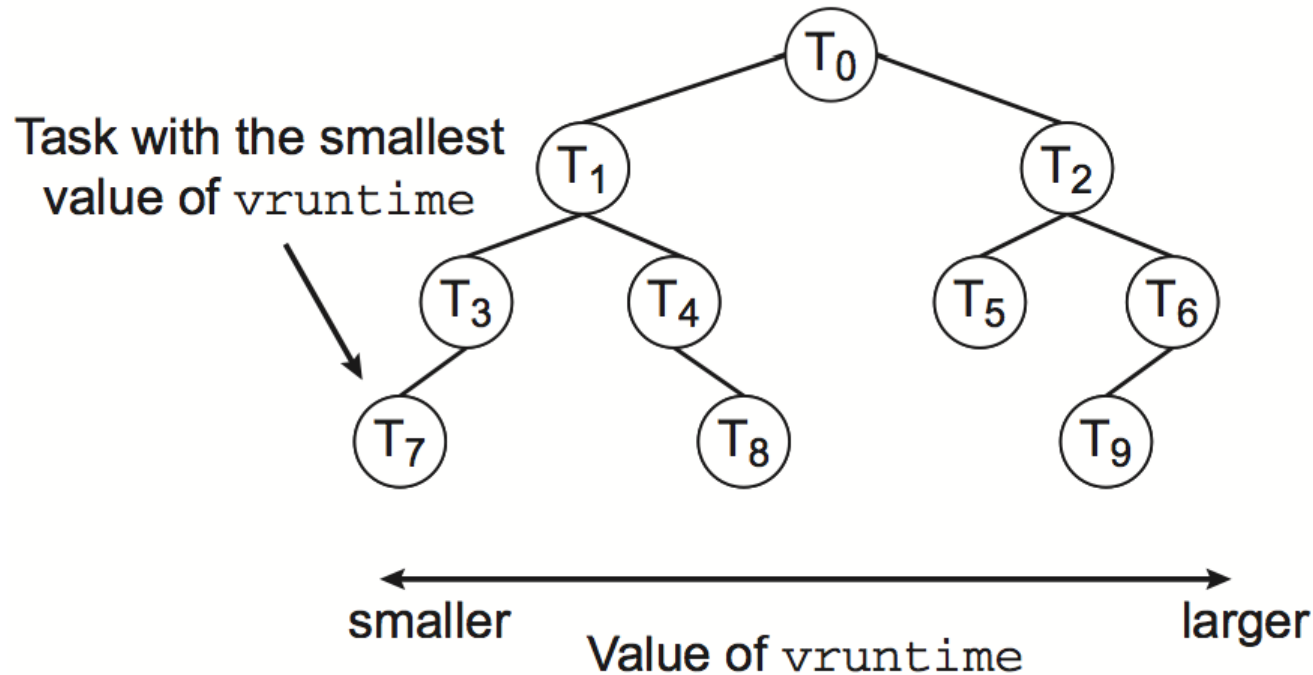
► Nice Value

- From -20 to +19
- Lower value is higher priority



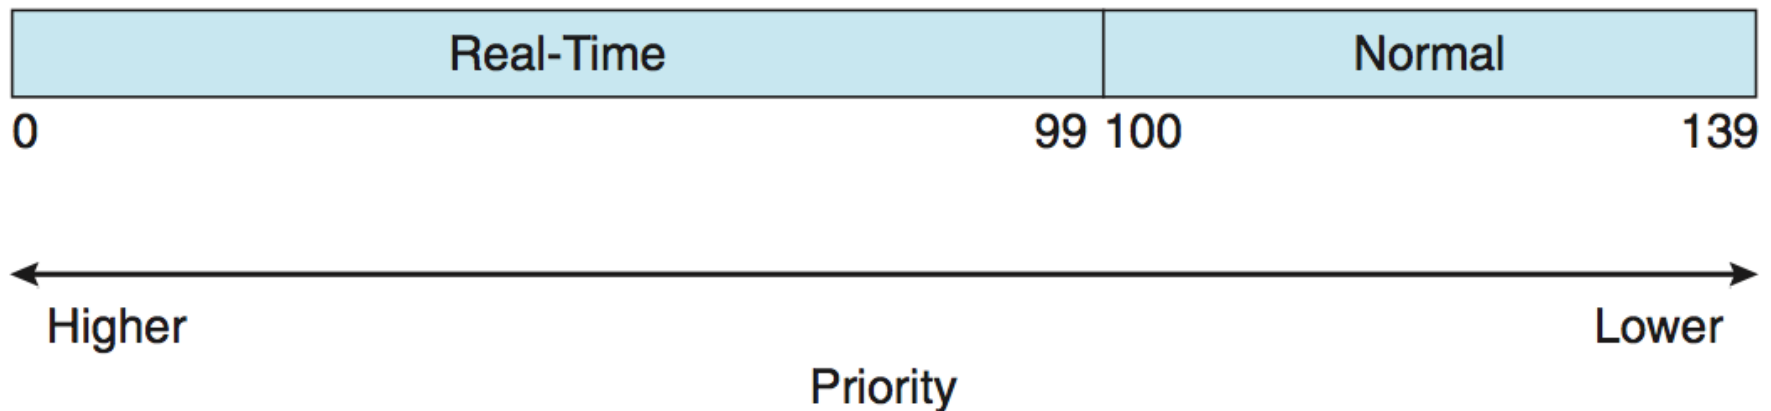
Operating System Examples – Linux in Version 2.6.23 + (2/3)

- ▶ A red-back tree is used to maintain the virtual run times of tasks



Operating System Examples – Linux in Version 2.6.23 + (3/3)

- ▶ Real-time scheduling according to POSIX
 - Real-time tasks have static priorities
- ▶ Real-time plus normal map into global priority scheme
- ▶ Nice value of -20 maps to global priority 100
- ▶ Nice value of +19 maps to priority 139



Operating System Examples – Windows Scheduling (1 / 3)

↓ A Typical Class

	Real-time	High	Above normal	Normal	Below normal	Idle priority
Time-critical	31	15	15	15	15	15
Highest	26	15	12	10	8	6
Above normal	25	14	11	9	7	5
Normal	24	13	10	8	6	4
Below normal	23	12	9	7	5	3
Lowest	22	11	8	6	4	2
Idle	16	1	1	1	1	1

Base
Priority →

Real-Time Class

Variable Class (1..15)



Operating System Examples – Windows Scheduling (2/3)

- ▶ Priority-Based Preemptive Scheduling
 - Priority Range: from 0 to 31
 - Variable class uses 1-15
 - Real-time class uses 16-31
 - Dispatcher: A process runs until
 - It is preempted by a higher-priority process
 - It terminates
 - Its time quantum ends
 - It calls a blocking system call
 - Idle thread
- ▶ A Queue per Priority Level



Operating System Examples – Windows Scheduling (3 / 3)

- ▶ Each thread has a base priority that represents a value in the priority range of its class
- ▶ Priority Changing
 - Increased after some waiting
 - Different amount for different I/O devices
 - Decreased after some computation
 - The priority is never lowered below the base priority
- ▶ Favor foreground processes
 - Each foreground task is given more time quantum (typically 3 times longer)



Scheduling Algorithm Evaluation

- ▶ A General Procedure
 - Select criteria that may include several measures, e.g., maximize CPU utilization while confining the maximum response time to 1 second
 - Evaluate various algorithms
- ▶ Evaluation Methods:
 - Deterministic modeling
 - Queuing models
 - Simulation
 - Implementation



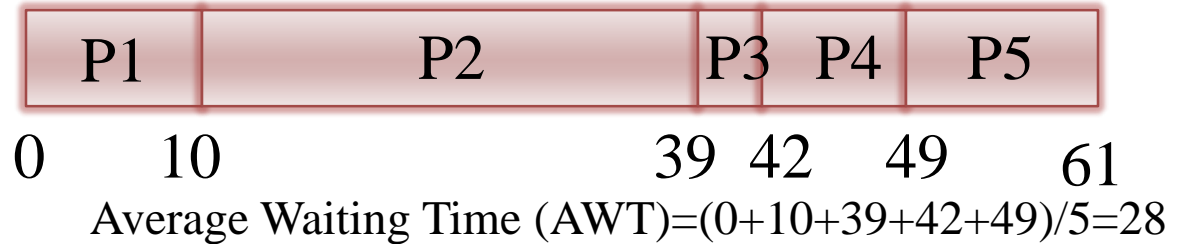
Deterministic Modeling

- ▶ A Typical Type of Analytic Evaluation
 - Take a particular predetermined workload and defines the performance of each algorithm for that workload
- ▶ Properties
 - Simple and fast
 - Through excessive executions of a number of examples, trends might be identified
 - But it needs exact numbers for inputs, and its answers only apply to those cases
 - Being too specific and requires too exact knowledge to be useful



Deterministic Modeling

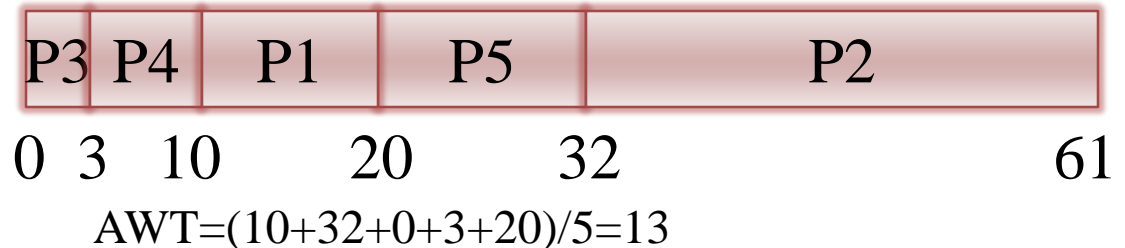
FCFS



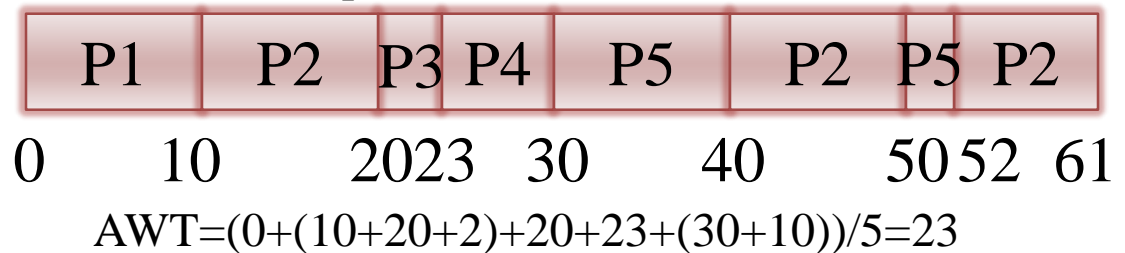
process CPU Burst time

P1	10
P2	29
P3	3
P4	7
P5	12

Nonpreemptive Shortest Job First



Round Robin (quantum = 10)



Queuing Models

- ▶ Motivation:
 - Workloads vary, and there is no static set of processes
- ▶ Models (~ Queuing-Network Analysis)
 - Workload:
 - Arrival rate: the distribution of times when processes arrive
 - The distributions of CPU & I/O bursts
 - Service rate



Simulation (1 / 2)

▶ Motivation:

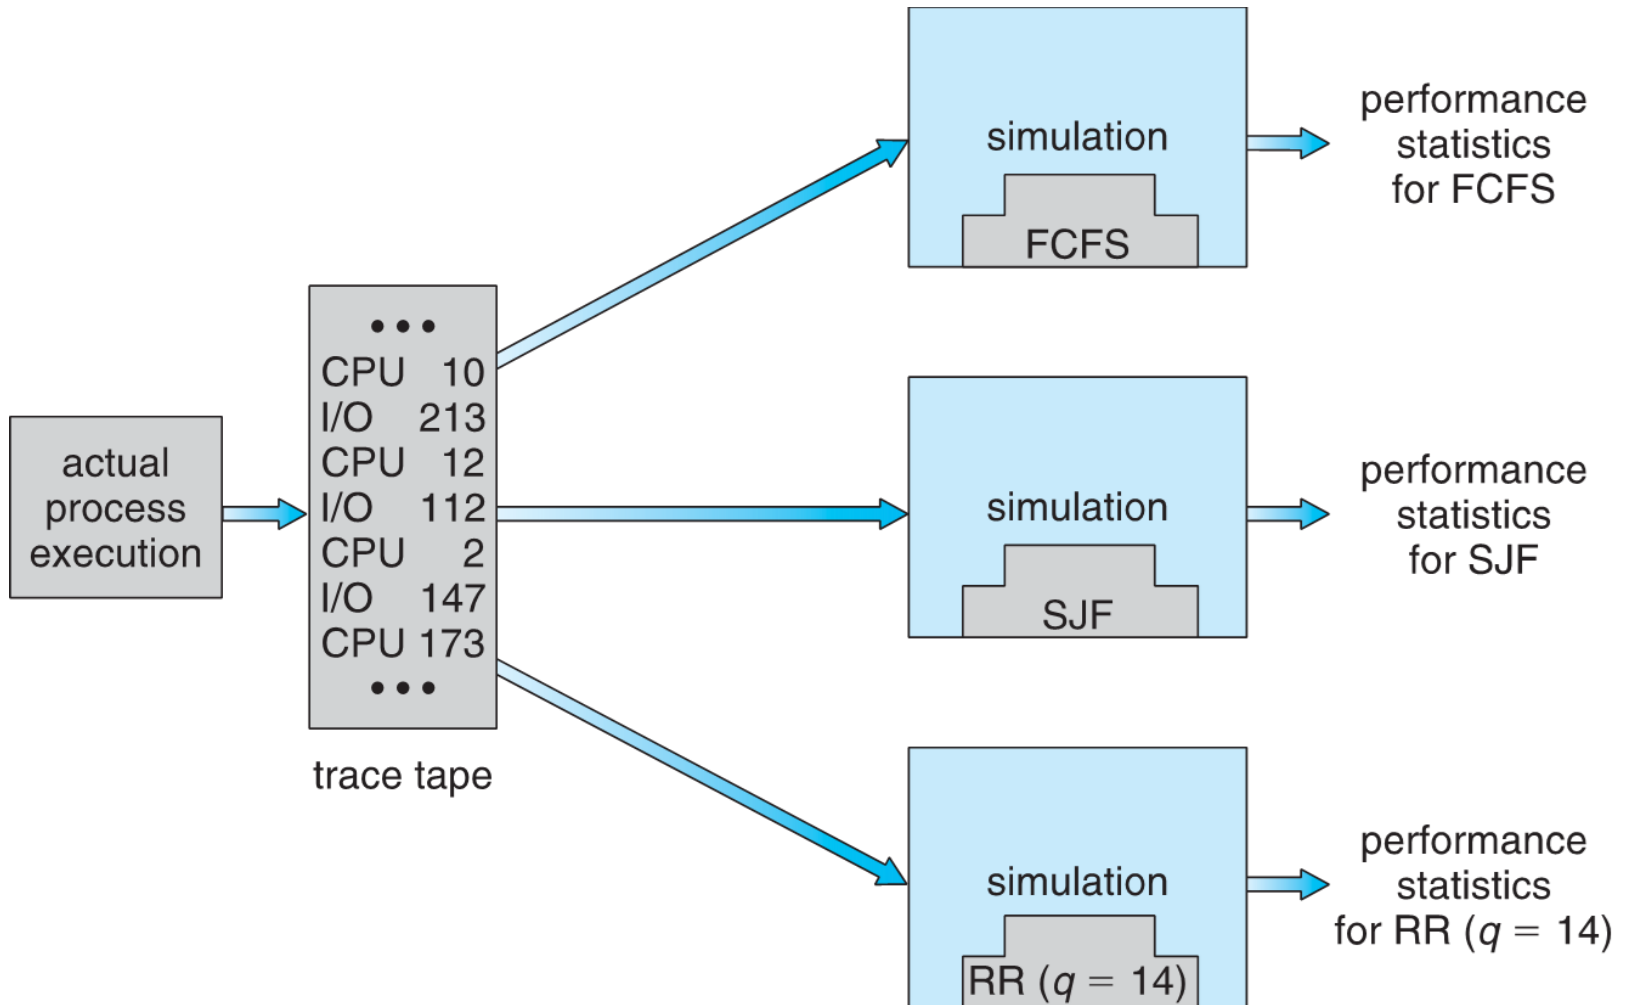
- Get a more accurate evaluation

▶ Procedures:

- Program a model of the computer system
- Drive the simulation with various data sets
 - Randomly generated according to some probability distributions
 - ➔ Inaccuracy occurs because of only the occurrence frequency of events. Miss the order & the relationships of events.
 - Trace tapes: monitor the real system & record the sequence of actual events.



Simulation (2/2)



Implementation

► Motivation:

- Get more accurate results than a simulation

► Procedure:

- Code scheduling algorithms
- Put them in the OS
- Evaluate the real behaviors

► Difficulties:

- Cost in coding algorithms and modifying the OS
- Reaction of users to a constantly changing the OS
- The environment in which algorithms are used will change

