



Operating System Concepts

Che-Wei Chang

chewei@mail.cgu.edu.tw

Department of Computer Science and Information Engineering, Chang Gung University

Contents

1. Introduction
2. System Structures
3. Process Concept
4. Multithreaded Programming
5. Process Scheduling
- ➔ 6. Synchronization
7. Deadlocks
8. Memory-Management Strategies
9. Virtual-Memory Management
10. File System
11. Implementing File Systems
12. Secondary-Storage Systems





Chapter 6.

Synchronization

Objectives

- ▶ To introduce the critical-section problem, whose solutions can be used to ensure the consistency of shared data
- ▶ To present both software and hardware solutions of the critical-section problem
- ▶ To examine several classical process-synchronization problems
- ▶ To explore several tools that are used to solve process synchronization problems



A Consumer-Producer Example

■ Producer

```
while (1) {  
    while (counter == BUFFER_SIZE)  
        ;  
    produce an item in nextp;  
    buffer[in] = nextp;  
    in = (in+1) % BUFFER_SIZE;  
    counter++;  
}
```

■ Consumer:

```
while (1) {  
    while (counter == 0)  
        ;  
    nextc = buffer[out];  
    out = (out + 1) % BUFFER_SIZE;  
    counter--;  
    consume an item in nextc;  
}
```



Race Condition

- ▶ One counter++ and one counter--

$r1 = \text{counter}$

$r2 = \text{counter}$

$r1 = r1 + 1$

$r2 = r2 - 1$

$\text{counter} = r1$

$\text{counter} = r2$

- ▶ Initially, let $\text{counter} = 5$

1. P: $r1 = \text{counter}$

2. P: $r1 = r1 + 1$

3. C: $r2 = \text{counter}$

4. C: $r2 = r2 - 1$  A Race Condition!

5. P: $\text{counter} = r1$

6. C: $\text{counter} = r2 = 4$

- ▶ The result can be 4, 5 or 6




Process Synchronization

▶ A Race Condition:

- A situation where the outcome of the execution depends on the particular order of process scheduling

▶ The Critical-Section Problem:

- Design a protocol that processes can use to cooperate
 - Each process has a segment of code, called a **critical section**, whose execution must be **mutually exclusive**
 - A general structure for the critical-section design



Solution of the Critical Section Problem

► Three Requirements

- Mutual Exclusion:

Only one process can be in its critical section

- Progress:

If no process is executing in its critical section and there exist some processes that wish to enter their critical section, then the selection of the processes that will enter the critical section next cannot be postponed indefinitely

- Bounded Waiting:

A waiting process only waits for a bounded number of processes to enter its critical section



Peterson's Solution (1 / 5)

▶ Notation

- Processes P_i and P_j

▶ Assumption

- Every basic machine-language instruction is atomic

▶ Algorithm 1

- Idea: Remember which process is allowed to enter its critical section. That is, P_i can enter its critical section if $\text{turn} = i$

do {

while (turn \neq i) ;

critical section

turn=j;

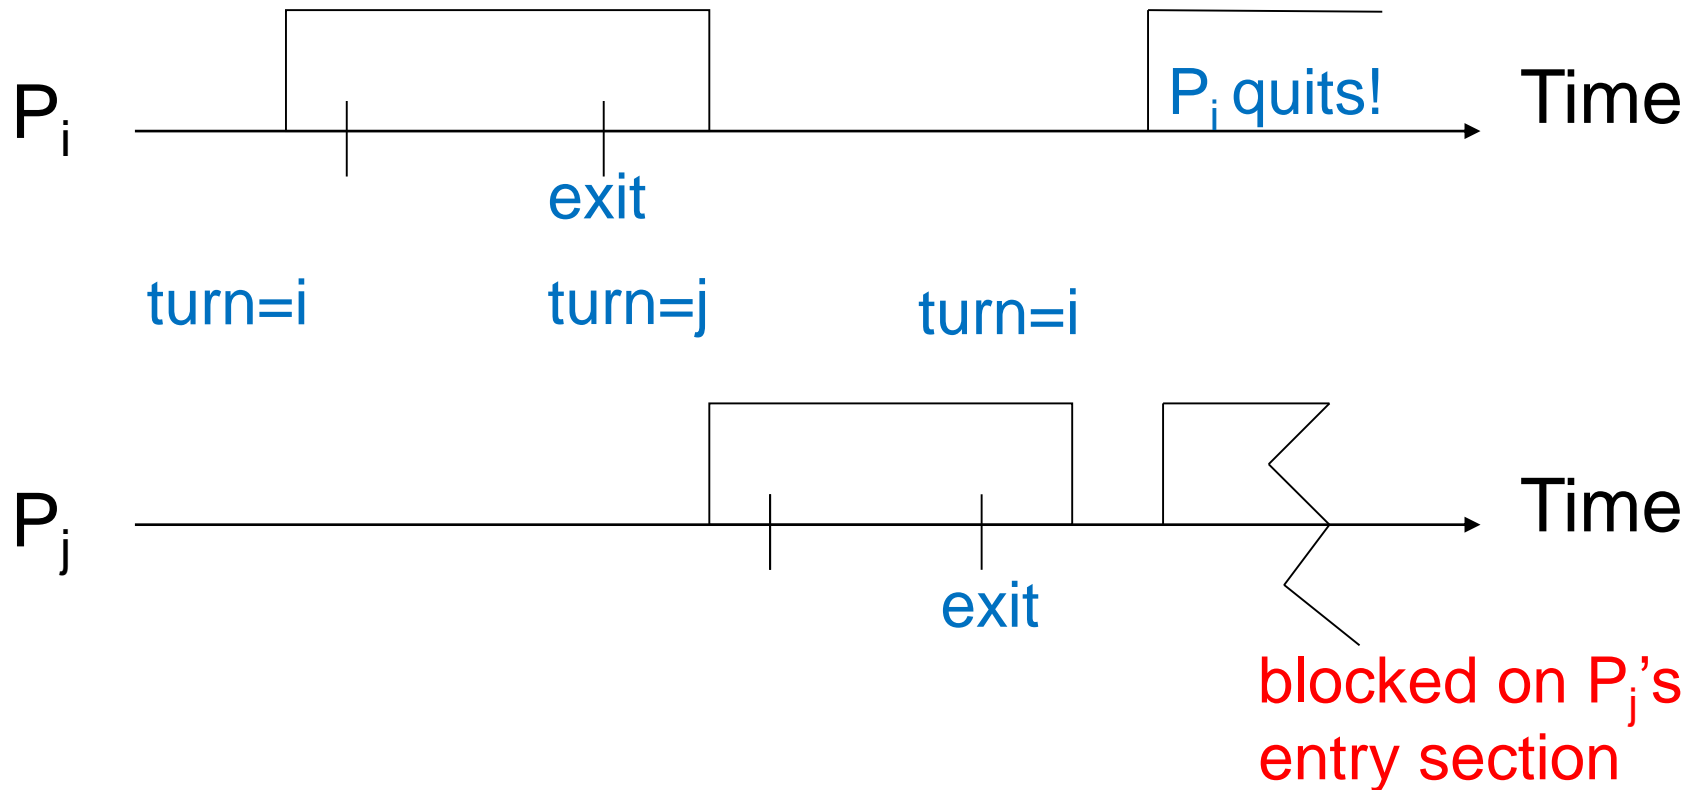
remainder section

} while (1);



Peterson's Solution (2 / 5)

- ▶ Algorithm 1 fails the progress requirement:



Peterson's Solution (3 / 5)

► Algorithm 2

- Idea: Remember the state of each process
- $\text{flag}[i] == \text{true} \rightarrow P_i$ is ready to enter its critical section
- Algorithm 2 fails the progress requirement when $\text{flag}[i] == \text{flag}[j] == \text{true}$

Initially, $\text{flag}[i] = \text{flag}[j] = \text{false}$

do {

```
flag[i]=true;
while (flag[j]) ;
```

critical section

```
flag[i]=false;
```

remainder section

} while (1);



Peterson's Solution (4/5)

► Algorithm 3

- Idea: Combine the ideas of Algorithms 1 and 2
- When $(\text{flag}[i] \ \&\& \ \text{turn}=i)$, P_j must wait
- Initially, $\text{flag}[i]=\text{flag}[j]=\text{false}$, and $\text{turn} = i$ or j

do {

```
flag[i]=true;
```

```
turn=j;
```

```
while (flag[j] && turn==j) ;
```

```
critical section
```

```
flag[i]=false;
```

```
remainder section
```

```
} while (1);
```



Peterson's Solution (5 / 5)

► Properties of Algorithm 3

◦ Mutual Exclusion

- The eventual value of *turn* determines which process enters the critical section

◦ Progress

- A process can only be stuck in the while loop, and the process which can keep it waiting must be in its critical sections

◦ Bounded Waiting

- Each process wait at most one entry by the other process



Peterson's Solution (4/5)

Process P_i 的程式碼:

do {

```
flag[i]=true;
```

```
turn=j;
```

```
while (flag[j] && turn==j) ;
```

critical section

```
flag[i]=false;
```

remainder section

```
} while (1);
```

Process P_j 的程式碼:

do {

```
flag[j]=true;
```

```
turn=i;
```

```
while (flag[i] && turn==i) ;
```

critical section

```
flag[j]=false;
```

remainder section

```
} while (1);
```



Brainstorming!

- Could we move $\text{turn}=i$; and $\text{turn}=j$; as follows:

Process P_i :

do {

$\text{flag}[i]=\text{true};$

$\text{while } (\text{flag}[j] \ \&\& \ \text{turn}==j) ;$

critical section

$\text{turn}=j;$

$\text{flag}[i]=\text{false};$

remainder section

} while (1);

Process P_j :

do {

$\text{flag}[j]=\text{true};$

$\text{while } (\text{flag}[i] \ \&\& \ \text{turn}==i) ;$

critical section

$\text{turn}=i;$

$\text{flag}[j]=\text{false};$

remainder section

} while (1);



Synchronization Hardware

▶ Motivation:

- Hardware features make programming easier and improve system efficiency

▶ Approach:

- Disable Interrupt → No Preemption
 - Infeasible in multiprocessor environments
 - Potential impacts on interrupt-driven system clocks
- Atomic Hardware Instructions
 - Test-and-set, Swap, etc.



Test and Set

Meaning

```
boolean TestAndSet(boolean *target) {  
    boolean rv = *target;  
    *target=true;  
    return rv;  
}
```

Usage

```
do {  
    while (TestAndSet(&lock)) ;  
    critical section  
    lock=false;  
    remainder section  
} while (1);
```



Swap

Meaning

```
void Swap(boolean *a, boolean *b) {  
    boolean temp = *a;  
    *a=*b;  
    *b=temp; }  

```

Usage

```
do {  
    key=true;  
    while (key == true)  
        Swap(&lock, &key);  
    critical section  
    lock=false;  
    remainder section  
} while (1);
```



Case Study of Test and Set

do {

▶ Problem

- n tasks want to access some share data

▶ Mutual Exclusion

- Pass if key== F or waiting[i]== F

▶ Progress

- Exit process sends a process in

▶ Bounded Waiting

- Wait at most n-1 times

```
waiting[i]=true;
key=true;
while (waiting[i] && key)
    key=TestAndSet(&lock);
waiting[i]=false;
```

critical section

```
j= (i+1) % n;
while( (j != i) && ( !waiting[j]) )
    j= (j+1) % n;
If (j=i) lock=false;
else waiting[j]=false;
```


remainder section

} while (1);



Mutex Locks (1 / 2)

- ▶ Previous solutions are complicated and generally inaccessible to application programmers
- ▶ OS designers build software tools to solve critical section problem
- ▶ Protect critical regions with it by first **acquire()** a lock then **release()** it
 - Boolean variable indicating if lock is available or not
- ▶ Calls to **acquire()** and **release()** must be atomic
 - Usually implemented via hardware atomic instructions
- ▶ But this solution requires **busy waiting**
 - This lock therefore called a **spinlock**



Mutex Locks (2 / 2)

Meaning

```
acquire() {  
    while (!available)  
        ; /* busy wait */  
    available = false;  
}  
release() {  
    available = true;  
}
```

Usage

```
do {  
    acquire(lock);  
    critical section  
    release(lock);  
    remainder section  
} while (true);
```



Semaphores (1 / 3)

- ▶ Motivation:
 - A high-level solution for more complex problems
- ▶ Semaphore
 - A variable **S** only accessible by two atomic operations:

```
wait(S) {      /* P */  
    while (S <= 0) ;  
    S--;  
}
```

```
signal(S) {    /* V */  
    S++;  
}
```



Semaphores (2 / 3)

► Critical Sections

```
do {  
    wait(S);  
    critical section  
    signal(S);  
    remainder section  
} while (1);
```

► Precedence Enforcement

```
P1:  
    S1;  
    signal(S);
```

```
P2:  
    wait(S);  
    S2;
```



Semaphores (3 / 3)

► Implementation

- Spinlock: A Busy-Waiting Semaphore
 - “while (S <= 0)” causes the wasting of CPU cycles!
 - Advantage:
 - When locks are held for a short time, spinlocks are useful since no context switching is involved.
- Semaphores with Blocked-Waiting
 - No busy waiting from the entry to the critical section!



Semaphores with Block Waiting

```
typedef struct{  
    int value;  
    struct process *list;  
} semaphore;
```

```
wait(semaphore *S) {  
    S->value--;  
    if (S->value < 0) {  
        add this process to S->list;  
        block();  
    }  
}
```

```
signal(semaphore *S) {  
    S->value++;  
    if (S->value <= 0) {  
        remove a process P from S->list;  
        wakeup(P);  
    }  
}
```



Deadlocks and Starvation

► Deadlock

- A set of processes is in a **deadlock** state when every process in the set is waiting for an event that can be caused only by another process in the set

P0: wait(S);	P1: wait(Q);
wait(Q);	wait(S);
...	...
signal(S);	signal(Q);
signal(Q);	signal(S);

► Starvation (or Indefinite Blocking)

- e.g., a LIFO (last-in, first-out) queue



Classical Problems of Synchronization

- ▶ Bounded-Buffer Problem
- ▶ Readers and Writers Problem
- ▶ Dining-Philosophers Problem



Bounded-Buffer Problem (1 / 2)

Producer:

do {

produce an item in nextp;

.....

wait(empty); /* control buffer availability */

wait(mutex); /* mutual exclusion */

.....


add nextp to buffer;

signal(mutex);

signal(full); /* increase item counts */

} while (1);

empty is
initialized to n



mutex is
initialized to 1



full is
initialized to 0



Bounded-Buffer Problem (2 / 2)


Consumer:

```
do {  
    wait(full); /* control buffer availability */  
    wait(mutex); /* mutual exclusion */  
    .....  
    remove an item from buffer to nextp;  
    .....  
    signal(mutex);  
    signal(empty); /* increase item counts */  
    consume nextp;  
} while (1);
```



Readers and Writers Problem (1 / 2)

- ▶ A data set is shared among a number of concurrent processes
 - Readers only read the data set; they do *not* perform any updates
 - Writers can both read and write
- ▶ Problem
 - Allow multiple readers to read at the same time
 - Only one single writer can access the shared data at the same time



Readers and Writers Problem (2 / 2)

```
semaphore wrt, mutex;  
    (initialized to 1);  
int readcount=0;
```

Writer:

```
➡ wait(wrt);  
.....  
writing is performed  
.....  
signal(wrt)
```



Reader:

```
➡ wait(mutex);  
  readcount++;  
  if (readcount == 1)  
      ➡ wait(wrt);  
  signal(mutex);  
  ... reading...  
➡ wait(mutex);  
  readcount--;  
  if (readcount == 0)  
      ➡ signal(wrt);  
  signal(mutex);
```



Dining-Philosopher Problem (1 / 3)

- ▶ Each philosopher must pick up one chopstick beside him/her at a time
- ▶ When two chopsticks are picked up, the philosopher can eat




Dining-Philosophers Problem (2 / 3)

```
semaphore chopstick[5];  
do {  
    wait(chopstick[i]);  
    wait(chopstick[(i + 1) % 5]);  
    ... eat ...  
    signal(chopstick[i]);  
    signal(chopstick[(i+1) % 5]);  
    ...think ...  
} while (1);
```



Dining-Philosophers Problem (3/3)

- ▶ This algorithm could create a deadlock
- ▶ Several possible remedies to the deadlock problem:
 - Allow at most four philosopher
 - Allow a philosopher to pick up chopsticks only if both are available
 - Asymmetric solution



Problems with Semaphores

- ▶ Incorrect use of semaphore operations:
 - signal (mutex) wait (mutex)
 - wait (mutex) ... wait (mutex)
 - Omitting of wait (mutex) or signal (mutex) (or both)
- ▶ Deadlock and starvation

➔ A high-level abstraction that provides a convenient and effective mechanism for process synchronization




Monitor (1 / 2)

► Components

- Variables
 - Monitor states
- Procedures
 - Only access local variables or formal parameters

monitor monitor-name


```
{  
    // shared variable declarations  
    procedure P1 (...) { .... }  
    procedure Pn (...) { ..... }  
    Initialization code (...) { ... }  
}
```



Monitor (2 / 2)

► Condition Variables

- `x.wait ()` – a process that invokes the operation is suspended until `x.signal ()`
- `x.signal ()` – resumes one of processes (if any) that invoked `x.wait ()`
 - If no `x.wait ()` on the variable, then it has no effect on the variable



Solution to Dining Philosophers (1 / 2)

monitor DiningPhilosophers

```
{
    enum { THINKING; HUNGRY, EATING) state [5] ;
    condition self [5];

    void pickup (int i) {
        state[i] = HUNGRY;
        test(i);
        if (state[i] != EATING) self [i].wait();
    }

    void putdown (int i) {
        state[i] = THINKING;
        // test left and right neighbors
        test((i + 4) % 5);
        test((i + 1) % 5);
    }
}
```

```
void test (int i) {
    if ( (state[(i + 4) % 5] != EATING) &&
        (state[i] == HUNGRY) &&
        (state[(i + 1) % 5] != EATING) ) {
        state[i] = EATING ;
        self[i].signal () ;
    }
}

initialization_code() {
    for (int i = 0; i < 5; i++)
        state[i] = THINKING;
}
}
```



Solution to Dining Philosophers (2/2)

- ▶ Each philosopher i invokes the operations `pickup()` and `putdown()` in the following sequence:

`DiningPhilosophers.pickup (i);`

Eat

`DiningPhilosophers.putdown (i);`

- ▶ No deadlock, but starvation is possible



Monitor Implementation Using Semaphores

- ▶ Semaphores
 - *mutex* – to protect the monitor
 - *next* – being initialized to zero, on which processes may suspend themselves
 - *next-count*
- ▶ For each external function F
 - `wait(mutex);`
 - ...
 - body of F ;*
 - ...
 - `if (next-count > 0)`
 - `signal(next);`
 - `else signal(mutex);`



Monitor Implementation Using Condition Variables

- For every condition x
 - A semaphore $x\text{-sem}$
 - An integer variable $x\text{-count}$
 - Implementation of $x.\text{wait}()$ and $x.\text{signal}()$:

- $x.\text{wait}()$

```
x-count++;  
if (next-count > 0)  
    signal(next);  
else  
    signal(mutex);  
wait(x-sem);  
x-count--;
```

- $x.\text{signal}()$

```
if (x-count > 0)  
{  
    next-count++;  
    signal(x-sem);  
    wait(next);  
    next-count--;  
}
```

* $x.\text{wait}()$ and $x.\text{signal}()$ are invoked within a monitor



Resuming Processes within a monitor

- ▶ How do we determine which of the suspended processes should be resumed next ?
 - FCFS ordering
 - Conditional-wait construct `x.wait(c);`
 - Monitor-scheduling algorithm
 - Built-in
 - or
 - User define

